

Use a Spark pool in Synapse to read a CSV file from ADLS and load the data into a Spark DataFrame.

The screenshot shows the 'Data' view in the Microsoft Azure Synapse Analytics portal. The left sidebar displays the workspace hierarchy: 'Workspace' > 'Linked' > 'Azure Data Lake Storage Gen2' > 'sparkpracmod (Primary - storagefo...)' > 'data (Primary)' > '(Attached Containers)'. The main pane shows a table of files in the 'data' container. A notification at the top right states 'Upload completed: 5 files uploaded to 'data''. The table lists five CSV files named 'orders1.csv' through 'orders5.csv', all with a content type of 'text/csv' and a size of 58 B. The 'Last Modified' timestamp for all files is '4/23/2026, 5:59:01 PM'. At the bottom, it says 'Showing 1 to 5 of 5 cached items'.

Name	Last Modified	Content Type	Size
orders1.csv	4/23/2026, 5:59:01 PM	text/csv	58 B
orders2.csv	4/23/2026, 5:59:00 PM	text/csv	58 B
orders3.csv	4/23/2026, 5:59:00 PM	text/csv	58 B
orders4.csv	4/23/2026, 5:59:00 PM	text/csv	58 B
orders5.csv	4/23/2026, 5:59:00 PM	text/csv	59 B

The screenshot shows the 'Develop' view in the Microsoft Azure Synapse Analytics portal. The left sidebar shows the workspace hierarchy: 'Workspace' > 'Linked' > 'Azure Data Lake Storage Gen2' > 'sparkpracmod (Primary - storagefo...)' > 'data (Primary)' > '(Attached Containers)'. The main pane shows a table of files in the 'data' container. A dropdown menu is open over the 'New' button, showing options: 'New SQL script', 'New notebook', 'New data flow', 'New integration dataset', 'Upload', 'Download', 'New folder', 'Select all', 'Rename', and 'More'. The table lists five CSV files named 'orders1.csv' through 'orders5.csv', all with a content type of 'text/csv' and a size of 58 B. The 'Last Modified' timestamp for all files is '4/23/2026, 5:59:01 PM'. At the bottom, it says 'Showing 1 to 5 of 5 cached items'.

Name	Last Modified	Content Type	Size
orders1.csv	4/23/2026, 5:59:01 PM	text/csv	58 B
orders2.csv	4/23/2026, 5:59:00 PM	text/csv	58 B
orders3.csv	4/23/2026, 5:59:00 PM	text/csv	58 B
orders4.csv	4/23/2026, 5:59:00 PM	text/csv	58 B
orders5.csv	4/23/2026, 5:59:00 PM	text/csv	59 B

Microsoft Azure | Synapse Analytics | sparkpracmod

Synapse live | Validate all | Publish all

Connector upgrade advisor

Analytics pools

SQL pools

Apache Spark pools

Data Explorer pools (preview)

External connections

Linked services

Microsoft Purview

Integration

Triggers

Integration runtimes

Security

Access control

Credentials

Managed private endpoints

Configurations + libraries

Workspace packages

Data flow libraries

Apache Spark configurations

Source control

Git configuration

Apache Spark pool

Apache Spark pools can be tuned to run different kinds of Apache S

+ New Refresh

Filter by name

Showing 0-0 of 0 item

Name

New Apache Spark pool

Basics | Additional settings | Tags | Review + create

Create an Synapse Analytics Apache Spark pool with your preferred configurations. Complete the Basics tab then go to Review + Create to provision with smart defaults, or visit each tab to customize.

Apache Spark pool details

Name your Apache Spark pool and choose its initial settings.

Apache Spark pool name * sparkpoolforlms

Node size family * Memory Optimized

Node size * Small (4 vCores / 32 GB)

Autoscale * ☒ Enabled ☐ Disabled

Number of nodes * 3

Estimated price ☐ Est. cost per hour 1.72 to 1.72 USD [View pricing details](#)

Dynamically allocate executors * ☐ Enabled ☒ Disabled

Activate Windows
Go to Settings to activate Windows.

Review + create | Next: Additional settings > | Cancel

Microsoft Azure | Synapse Analytics | sparkpracmod

Synapse live | Validate all | Publish all

Connector upgrade advisor

Analytics pools

SQL pools

Apache Spark pools

Data Explorer pools (preview)

External connections

Linked services

Microsoft Purview

Integration

Triggers

Integration runtimes

Security

Access control

Credentials

Managed private endpoints

Configurations + libraries

Workspace packages

Data flow libraries

Apache Spark configurations

Source control

Git configuration

Apache Spark pool

Apache Spark pools can be tuned to run different kinds of Apache S

+ New Refresh

Filter by name

Showing 0-0 of 0 item

Name

New Apache Spark pool

Basics | Additional settings | Tags | Review + create

Validation succeeded.

Product details

Azure Synapse Analytics Apache Spark pool by Microsoft ☐ Est. cost per hour 1.72 to 1.72 USD [View pricing details](#)

Terms

By clicking "Create", I (a) agree to the legal terms and privacy statement(s) associated with the Marketplace offering(s) listed above; (b) authorize Microsoft to bill my current payment method for the fees associated with the offering(s), with the same billing frequency as my Azure subscription; and (c) agree that Microsoft may share my contact, usage and transactional information with the provider(s) of the offering(s) for support, billing and other transactional activities. Microsoft does not provide rights for third-party offerings. For additional details see [Azure Marketplace Terms](#)

Basics

Subscription Azure for Students

Resource group moduleonelms

Apache Spark pool name sparkpoolforlms

Node size family Memory Optimized

Node size Small (4 vCores / 32 GB)

Autoscale Enabled

Number of nodes(min/max) 3 to 3 nodes

Additional settings

Automatic pausing Enabled

Number of minutes idle 15

Dynamically allocate executors Disabled

Activate Windows
Go to Settings to activate Windows.

Create | < Previous | Download template for automation | Cancel

Microsoft Azure | Synapse Analytics | sparkpracmod

Search

kolletushar@outlook.com
DEFAULT DIRECTORY

Synapse live

Validate all

Publish all

Workspace

Linked

Filter resources by name

Azure Data Lake Storage Gen2

2

sparkpracmod (Primary - storagefo...

data (Primary)

(Attached Containers)

data

Notebook 1

Other users in your workspace may have access to modify this item. Do not use this item unless you trust all users who may have access to the workspace.

Run all

Undo

Publish

Outline

Attach to sparkpoolforins

Language PySpark (Python)

Variables

Ready

```
1 df = spark.read.format("csv") \
2   .option("header", "true") \
3   .option("inferSchema", "true") \
4   .load("abfss://data@storageforlmsdemo.dfs.core.windows.net/orders1.csv")
5
6 df.show()
```

[?] 3 min 59 sec - Apache Spark session started in 3 min 40 sec 291 ms. Command executed in 18 sec 898 ms by kolletushar on 6/18/38 PM, 4/23/26

Job execution Succeeded

Spark 2 executors 8 cores

View in monitoring

Open Spark UI

OrderID	CustomerID	Amount
101	1	250
102	2	450
103	3	300

Properties

General

Related (0)

Name *
Notebook 1

Description

Type
.ipynb notebook

Size
2,075 bytes

Notebook settings

☒ Include cell output when saving

☐ Enable unpublished notebook reference

Session
[Configure session](#)

Activate Windows
Go to Settings to activate Windows.